

IT2201 / IT2601 / IT2564 / IT2621 / IT2521 / IT2323

Database Management Systems



## Unit 10

### Transaction Management

1

## Unit Objectives

---

- ▣ Understand Transaction Support and their properties
- ▣ Understand Concurrency Control Services.
- ▣ The ability to identify problems caused by concurrency and how to prevent it.
- ▣ Understand database recovery, identify areas of failure and how to prevent them.

2

## Transaction Support

### □ What is a Transaction

- Action, or series of actions, carried out by user or application, which accesses or changes contents of database.

### □ A transaction is a logical unit of work on the database.

### □ A transaction transforms database from one consistent state to another.

### □ A transaction can have one of two outcomes:

- commit
- rollback

3

## Example of a Transaction

### □ Transaction (a)

- To update the salary of a particular member of staff given the staff number, x.

### □ Transaction (b)

- To delete the member of staff with a given staff number x, and to reassign the properties managed by this staff to a different member of staff.

```
read(staffNo = x, salary)
salary = salary * 1.1
write(staffNo = x, new_salary)
```

(a)

```
delete(staffNo = x)
for all PropertyForRent records, pno
begin
  read(propertyNo = pno, staffNo)
  if (staffNo = x) then
    begin
      staffNo = newStaffNo
      write(propertyNo = pno, staffNo)
    end
  end
end
```

(b)

4

## Transaction Support

- ❑ The DBMS provides the following keywords to allow the user to indicate the boundaries of a transaction:
  - BEGIN TRANSACTION
  - COMMIT
  - ROLLBACK
- ❑ If these delimiters are not used,
  - the entire program is regarded as a single transaction
  - DBMS automatically performing a COMMIT when the program terminates correctly and a ROLLBACK if it does not.

5

## Properties of Transactions

- ❑ The four basic properties that all transactions should possess are commonly known as the **ACID Properties**:
  - Atomicity
    - ❑ The 'all or nothing' property. A transaction is an indivisible unit that is either performed in its entirety or is not performed at all.
  - Consistency
    - ❑ A transaction must transform database from one consistent state to another.
  - Isolation
    - ❑ Transactions execute independently of one another. That is, the partial effects of incomplete transactions should not be visible to other transactions.
  - Durability
    - ❑ Effects of a committed transaction are permanent and must not be lost because of later failure.

6

## Review Question 1

- ▣ Explain the “ACID” characteristics of the transaction and how it helps maintain data integrity.

7

## Concurrency Control

- ▣ What is concurrency control
  - The process of managing simultaneous operations on the database without having them interfere with one another.
- ▣ Problems that can arise with concurrent access are :
  - Lost update problem
  - Uncommitted dependency (or dirty read) problem
  - Inconsistent analysis problem

8

## Concurrency Control

### ❑ Lost update problem

- This problem occurs when an successfully completed update operation by one user is overridden by another user.

### ❑ Uncommitted dependency problem

- This problem occurs when one transaction is allowed to see the intermediate results of another transaction before it has committed.

### ❑ Inconsistent analysis problem

- This problem occurs when a transaction reads several values from the database but a second transaction updates some of them during execution of the first, and causing the prior transaction to obtain inaccurate results.

9

## Lost Update Problem

Time	T <sub>1</sub>	T <sub>2</sub>	bal <sub>x</sub>
t <sub>1</sub>		begin_transaction	100
t <sub>2</sub>	begin_transaction	read(bal <sub>x</sub> )	100
t <sub>3</sub>	read(bal <sub>x</sub> )	bal <sub>x</sub> = bal <sub>x</sub> + 100	100
t <sub>4</sub>	bal <sub>x</sub> = bal <sub>x</sub> - 10	write(bal <sub>x</sub> )	200
t <sub>5</sub>	write(bal <sub>x</sub> )	commit	90
t <sub>6</sub>	commit		90

10

## Uncommitted Dependency Problem (or dirty read)

Time	T <sub>3</sub>	T <sub>4</sub>	bal <sub>x</sub>
t <sub>1</sub>		begin_transaction	100
t <sub>2</sub>		read(bal <sub>x</sub> )	100
t <sub>3</sub>		bal <sub>x</sub> = bal <sub>x</sub> + 100	100
t <sub>4</sub>	begin_transaction	write(bal <sub>x</sub> )	200
t <sub>5</sub>	read(bal <sub>x</sub> )	:	200
t <sub>6</sub>	bal <sub>x</sub> = bal <sub>x</sub> - 10	rollback	100
t <sub>7</sub>	write(bal <sub>x</sub> )		190
t <sub>8</sub>	commit		190

11

## Inconsistent Analysis Problem

Time	T <sub>5</sub>	T <sub>6</sub>	bal <sub>x</sub>	bal <sub>y</sub>	bal <sub>z</sub>	sum
t <sub>1</sub>		begin_transaction	100	50	25	
t <sub>2</sub>	begin_transaction	sum = 0	100	50	25	0
t <sub>3</sub>	read(bal <sub>x</sub> )	read(bal <sub>x</sub> )	100	50	25	0
t <sub>4</sub>	bal <sub>x</sub> = bal <sub>x</sub> - 10	sum = sum + bal <sub>x</sub>	100	50	25	100
t <sub>5</sub>	write(bal <sub>x</sub> )	read(bal <sub>y</sub> )	90	50	25	100
t <sub>6</sub>	read(bal <sub>z</sub> )	sum = sum + bal <sub>y</sub>	90	50	25	150
t <sub>7</sub>	bal <sub>z</sub> = bal <sub>z</sub> + 10		90	50	25	150
t <sub>8</sub>	write(bal <sub>z</sub> )		90	50	35	150
t <sub>9</sub>	commit	read(bal <sub>z</sub> )	90	50	35	150
t <sub>10</sub>		sum = sum + bal <sub>z</sub>	90	50	35	185
t <sub>11</sub>		commit	90	50	35	185

12

## Concurrency Control Technique

### □ Locking

- A procedure used to control concurrent access to data. When one transaction is accessing the database, a lock may deny access to other transactions to prevent incorrect results.

### □ Generally

- A transaction must claim a **shared** (read) or **exclusive** (write) lock on a data item before the corresponding database read or write operation.
- The lock prevents another transaction from modifying item or even reading it, in the case of an exclusive lock.

13

## Basic Rules for Locking

- If a transaction has a **shared lock** on a data item, it can read the item but not update it.
- If a transaction has an **exclusive lock** on a data item, it can both read and update the item.
- Since reads operations cannot conflict, so more than one transaction can hold shared locks simultaneously on same item.
- An exclusive lock gives a transaction exclusive access to that item.
- Some systems allow transaction to upgrade shared lock to a exclusive lock, or downgrade exclusive lock to a shared lock.

14

## Preventing Lost Update Problem

Time	T <sub>1</sub>	T <sub>2</sub>	bal <sub>x</sub>
t <sub>1</sub>		begin_transaction	100
t <sub>2</sub>	begin_transaction	write_lock(bal <sub>x</sub> )	100
t <sub>3</sub>	write_lock(bal <sub>x</sub> )	read(bal <sub>x</sub> )	100
t <sub>4</sub>	WAIT	bal <sub>x</sub> = bal <sub>x</sub> + 100	100
t <sub>5</sub>	WAIT	write(bal <sub>x</sub> )	200
t <sub>6</sub>	WAIT	commit/unlock(bal <sub>x</sub> )	200
t <sub>7</sub>	read(bal <sub>x</sub> )		200
t <sub>8</sub>	bal <sub>x</sub> = bal <sub>x</sub> - 10		200
t <sub>9</sub>	write(bal <sub>x</sub> )		190
t <sub>10</sub>	commit/unlock(bal <sub>x</sub> )		190

15

## Review Question 2

- ▣ What are the problems identified in concurrency control?
- ▣ Name the technique used to overcome the problems mentioned.
- ▣ What are the 2 commonly used locks?

16



## Database Recovery

### □ What is database recovery


- The process of restoring the database to a correct state in the event of a failure.

### □ Types of failure

- System failures (soft crash) - E.g. power failure
- Hardware failures (hard crash) - E.g. disk head crash
- Logical errors - E.g. integer overflow or division by zero


17

## Process of a database update operation

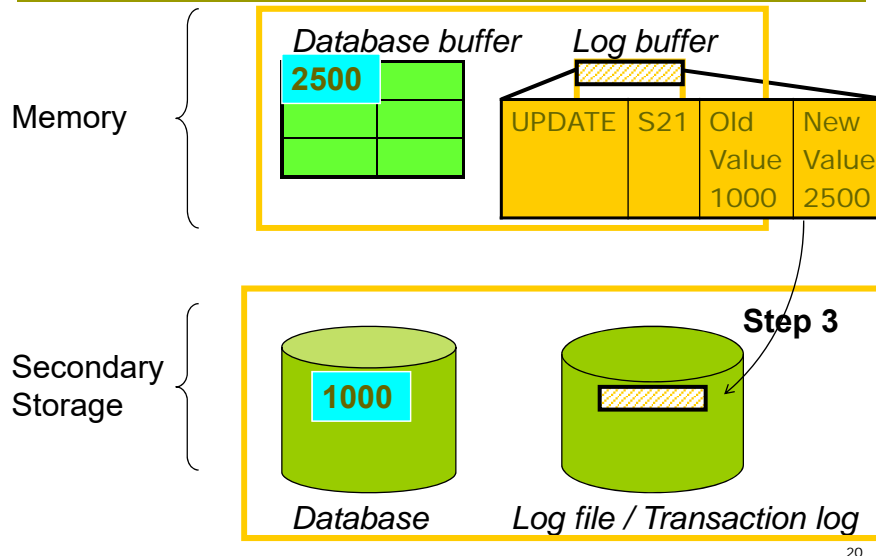


18

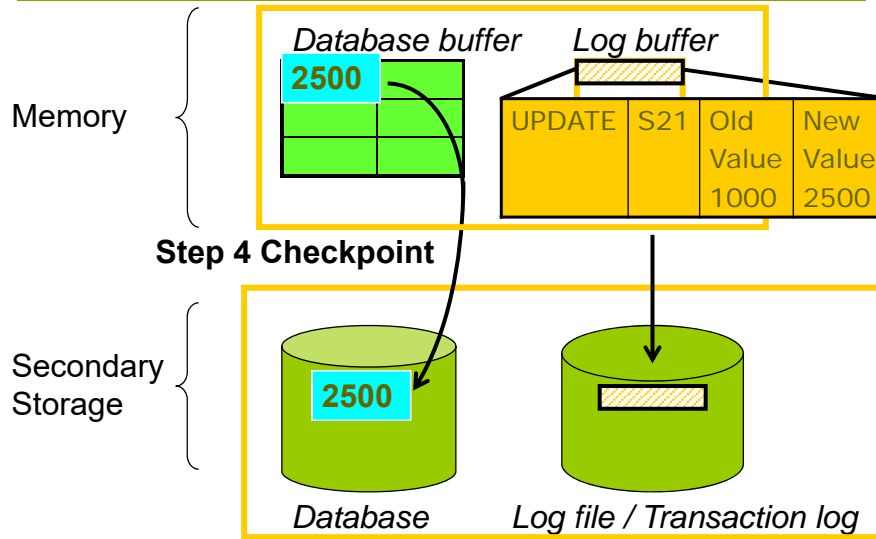
## Process of a database update operation



## Process of a database update operation



## Process of a database update operation



## Log File

- A log file is a special file maintained by the DBMS, which contains information about all updates to database.

- E.g. of a update record in the log file :

Tid	Time	Operation	Object	Before Image	After Image	pPtr	nPtr
T1	10:12	START				0	2
T1	10:13	UPDATE	STAFF SL21	(old value)	(new value)	1	8

- Note that Before-image and After-image of data item are stored in the log file.
- The log is used for purposes such as
  - Recovery
  - Performance monitoring
  - Auditing

22

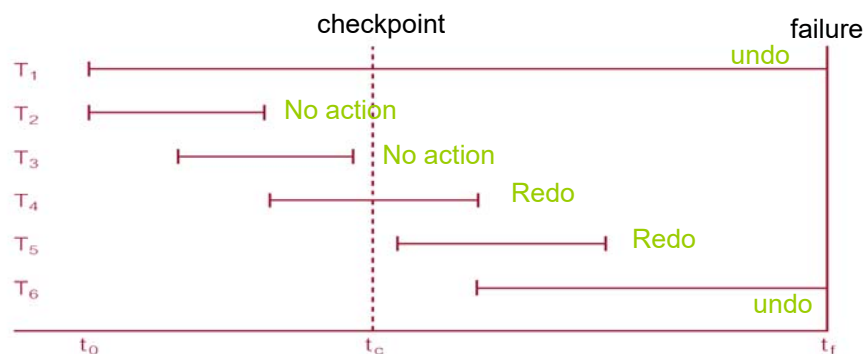
## Checkpointing

### □ Checkpoint

- It is the point of synchronization between the database and the log file. All buffers (including the database buffers and log buffer) are force-written to secondary storage.
- System automatically 'takes a checkpoint' at some prescribed intervals.
- At checkpoint, a checkpoint record is also written to the log file, and it contains all active transactions at the time of the checkpoint.

23

## Transactions and Recovery



How the recovery Manager perform recovery for the various transactions when there is a failure?

24

## Recovery Facilities

- ▣ A DBMS should provide the following facilities to assist with recovery:
  - Backup mechanism
    - ▣ makes periodic backup copies of database.
  - Logging facilities
    - ▣ keep track of current state of transactions and database changes.
  - Checkpoint facility
    - ▣ enables updates to the database that are in progress to be made permanent.
  - Recovery manager
    - ▣ allows the system to restore the database to a consistent state following a failure.

25

## Review questions 3

- ▣ List 3 types of database failures
- ▣ What are the uses for a log file being implemented in a Database
- ▣ What is the relationship of a log file and check point?

26

## Reference Materials

---

1. Database Systems, Connolly, Ch 22
2. Database Design, Application Development & Administration, Michael V. Mannino, Ch 13

27